

Step 1: Create and Activate Virtual Environment (if not already)

1. sudo apt update -y

```
adypuatharv@atharva-VirtualBox:~$ sudo apt upgrade -y
[sudo] password for adypuatharv:
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
Calculating upgrade... Done
The following packages were automatically installed and are no longer required:
  libgl1-amber-dri libglapi-mesa libllvm17t64 python3-netifaces
Use 'sudo apt autoremove' to remove them.
The following NEW packages will be installed:
  libllvm20 linux-headers-6.14.0-35-generic linux-hwe-6.14-headers-6.14.0-35 linux-hwe-6
  linux-image-6.14.0-35-generic linux-modules-6.14.0-35-generic linux-modules-extra-6.14
  linux-tools-6.14.0-35-generic mesa-libgallium
The following packages have been kept back:
  libgl1-amber-dri libglapi-mesa
```

2. sudo apt install python3-pip python3.12-venv -y

```
atharv@Atharv:~$ sudo apt install python3-pip python3-venv -y
Reading package lists... Done
Building dependency tree... Done
Reading state information... Done
The following packages were automatically installed and are no longer required:
  libdrm-nouveau2 libdrm-radeon1 libgl1-amber-dri libglapi-mesa libllvm17t64 libxcb-dri2-0
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed:
  build-essential bzip2 cpp cpp-13 cpp-13-x86-64-linux-gnu cpp-x86-64-linux-gnu dpkg-dev fakeroot g++ g++-13 g++-13-x86-64-linux-gnu g++-x86-64-linux-gnu
  gcc gcc-13 gcc-13-base gcc-13-x86-64-linux-gnu gcc-x86-64-linux-gnu javascript-common libalgorithm-diff-perl libalgorithm-diff-xs-perl
  libatomic1 libatomic1 libc-dev-bin libc-devtools libc6-dev libcc1-8 libcrypt-dev libde265-0 libdpkg-perl libexpat1-dev
  libfakeroot libfile-fcntllock-perl libgcc-13-dev libgd3 libheif-plugin-aomdec libheif-plugin-aomenc libheif-plugin-libde265 libheif1 libhwasan0 libisl23
  libitm1 libjs-jquery libjs-sphinxdoc libjsunderscore liblsan0 libmpc3 libpython3-dev libpython3.12-dev libquadmath0 libstdc++-13-dev libtsan2 libubsan1
  linux-libc-dev lto-disabled-list make manpages-dev python3-pip-whl python3-setuptools-whl python3-wheel python3.12-dev python3.12-venv
  rpcsvc-proto zlib1g-dev
Suggested packages:
  bzip2-doc cpp-doc gcc-13-locales cpp-13-doc debian-keyring g++-multilib g++-13-multilib gcc-13-doc gcc-multilib autoconf automake libtool flex bison gdb
  gcc-doc gcc-13-multilib gdb-x86-64-linux-gnu apache2 | lighttpd | httpd glibc-doc bzr libgd-tools libheif-plugin-x265 libheif-plugin-ffmpegdec
  libheif-plugin-jpegdec libheif-plugin-jpegenc libheif-plugin-j2kdec libheif-plugin-j2kenc libheif-plugin-rav1e libheif-plugin-svtenc libstdc++-13-doc
  make-doc
```

The command I ran earlier had already been executed, which is why it's showing only partially.

Normally, it takes some time to complete.

3. python3 -m venv airflow_env

```
atharv@Atharv:~/airflow$ python3 -m venv airflow_env
atharv@Atharv:~/airflow$ |
```

source airflow_env/bin/activate

```
atharv@Atharv:~/airflow$ source airflow_env/bin/activate
(airflow_env) atharv@Atharv:~/airflow$ |
```

```
atharv@Atharv:~$ source airflow_env/bin/activate
(airflow_env) atharv@Atharv:~$ pip install --upgrade pip
Requirement already satisfied: pip in ./airflow_env/lib/python3.12/site-packages (24.0)
Collecting pip
  Downloading pip-25.3-py3-none-any.whl.metadata (4.7 kB)
  Downloading pip-25.3-py3-none-any.whl (1.8 MB) 1.8/1.8 MB 475.7 kB/s eta 0:00:00
Installing collected packages: pip
  Attempting uninstall: pip
    Found existing installation: pip 24.0
    Uninstalling pip-24.0:
      Successfully uninstalled pip-24.0
Successfully installed pip-25.3
(airflow_env) atharv@Atharv:~$ pip install "apache-airflow==2.9.2" --constraint "https://raw.githubusercontent.com/apache/airflow/constraints-2.9.2/constraints-3.10.txt"
Collecting apache-airflow==2.9.2
  Downloading apache_airflow-2.9.2-py3-none-any.whl.metadata (43 kB)
Collecting alembic<2.0,>=1.13.1 (from apache-airflow==2.9.2)
  Downloading alembic-1.13.1-py3-none-any.whl.metadata (7.4 kB)
Collecting argcomplete>=1.10 (from apache-airflow==2.9.2)
  Downloading argcomplete-3.3.0-py3-none-any.whl.metadata (16 kB)
Collecting asgiref (from apache-airflow==2.9.2)
  Using cached asgiref-3.8.1-py3-none-any.whl.metadata (9.3 kB)
```

Step 2: Install Apache Airflow (Stable Version with Dependencies)

You need to add the ~/airflow_env/bin to your PATH permanently.

```
echo 'export PATH="$HOME/airflow_env/bin:$PATH"' >> ~/.bashrc
source ~/.bashrc
```

Step 4: Set Airflow Home Directory (Optional but Recommended)

By default, Airflow creates files in ~/airflow. You can choose a custom directory.

```
echo 'export AIRFLOW_HOME=~/airflow' >> ~/.bashrc
source ~/.bashrc
```

```
(airflow_env) atharv@Atharv:~/airflow$ export AIRFLOW_HOME=~/airflow
export AIRFLOW__CORE__LOAD_EXAMPLES=False
(airflow_env) atharv@Atharv:~/airflow$ |
```

Step 5: Initialize Airflow Database

```
airflow db init
```

All the commands are run in a single screenshot

```

airflow 0.2.0 airflow 2.0.0 aiida 2020.6.1 avrocon 0.50.0 avroloop 0.22.1 was
(airflow_env) atharv@Atharv:~/airflow$ airflow db init
Usage: airflow db [-h] COMMAND ...

Database operations

Positional Arguments:
  COMMAND
    check           Check if the database can be reached
    check-migrations  Check if migration have finished
    clean            Purge old records in metastore tables
    downgrade        Downgrade the schema of the metadata database.
    drop-archived    Drop archived tables created through the db clean com
    export-archived   Export archived data from the archive tables
    migrate          Migrates the metadata database to the latest version
    reset            Burn down and rebuild the metadata database
    shell             Runs a shell to access the database

Options:
  -h, --help      show this help message and exit

airflow db command error: argument COMMAND: invalid choice: 'init' (choose
from 'check', 'check-migrations', 'clean', 'downgrade', 'drop-archived',
'migrate', 'reset', 'shell'), see help above.
(airflow_env) atharv@Atharv:~/airflow$ |

```

Step 6: Create Airflow Admin User

```

airflow users create \
--username admin \
--firstname Admin \
--lastname User \
--role Admin \
--email admin@example.com

```

It will ask for password → set a secure one.

```

(airflow_env) atharv@Atharv:~/airflow$ airflow users create \
--username admin \
--firstname Atharv \
--lastname Jagtap \
--role Admin \
--email atharv@example.com
Usage: airflow [-h] GROUP_OR_COMMAND ...

Positional Arguments:
  GROUP_OR_COMMAND

  Groups
    assets           Manage assets
    backfill         Manage backfills
    config           View configuration
    connections     Manage connections
    dags             Manage DAGs
    db               Database operations
    db-manager       Manage externally connected database managers
    jobs             Manage jobs
    pools            Manage pools
    providers        Display providers

```

The above screenshot shows that the database exit because I have already runned this code first now when we run the second time it shows admin exist in the db

Step 7: Create folders for DAGs

```
mkdir -p ~/airflow/dags
```

Airflow will automatically look for DAGs in ~/airflow/dags.

Step 8 : Create a simple DAG

1. Open a text editor and create a file: nano ~/airflow/dags/etl_dag.py

2. Paste the following **example DAG code**:

```
from datetime import datetime
from airflow import DAG
from airflow.operators.python import PythonOperator
# Example ETL functions
def extract():
    print("Extracting data...")
def transform():
    print("Transforming data...")
def load():
    print("Loading data...")
# Define DAG
with DAG(
    'etl_dag',
    start_date=datetime(2025, 11, 7),
    schedule_interval='@daily',
    catchup=False,
    description='A simple ETL DAG with dependencies'
) as dag:
    # Define tasks
    extract_task = PythonOperator(
        task_id='extract',
        python_callable=extract
    )
    transform_task = PythonOperator(
        task_id='transform',
```

```
    python_callable=transform
)
load_task = PythonOperator(
    task_id='load',
    python_callable=load
)
# Define the task dependencies (graph)
extract_task >> transform_task >> load_task
```

Save and exit (Ctrl+O, Enter, Ctrl+X).

Step : Start Airflow services

1. Start the webserver:

```
airflow webserver --port 8080
```

By default, Airflow UI will be at <http://localhost:8080>.

```
---- /| [---()----- /--/----  
---- /| / /----/ /--/----/ \_ | /| / /  
---- /| / /----/ /----/ / / /----/ | / /  
Running the Gunicorn Server with:  
Workers: 4 sync  
Host: 0.0.0.0:8080  
Timeout: 120  
Logfiles: - -  
Access Logformat:  
===== /home/ubuntu/airflow_env/lib/python3.12/site-packages/flask_limiter/extension.py:333 UserWarning: Using the in-memory storage for tracking rate limits as no storage was explicitly specified. This is not recommended for production use. See: https://flask-limiter.readthedocs.io#configuring-a-storage-backend for documentation about configuring the storage backend.  
[2025-11-08 13:00:07 +0000] [717] [INFO] Starting gunicorn 22.0.0  
[2025-11-08 13:00:07 +0000] [717] [INFO] Listening at: http://0.0.0.0:8080 (717)  
[2025-11-08 13:00:07 +0000] [717] [INFO] Using worker: sync  
[2025-11-08 13:00:07 +0000] [719] [INFO] Booting worker with pid: 719  
[2025-11-08 13:00:07 +0000] [720] [INFO] Booting worker with pid: 720  
[2025-11-08 13:00:08 +0000] [721] [INFO] Booting worker with pid: 721  
[2025-11-08 13:00:08 +0000] [722] [INFO] Booting worker with pid: 722
```

2. In another terminal, activate the virtual environment and start the scheduler:

```
source airflow_env/bin/activate
```

airflow scheduler

Login in using username and password

The above dag with name etl_dag is the dag which we created now using the code

The below Screenshot represents the detail of the dag

The below screenshot represents the graph of the dag



